

Music Data Analysis Project

Project Description:

A leading music-catering company is planning to analyze large amount of data received from varieties of sources, namely mobile app and website to track the behavior of users, classify users, calculate royalties associated with the song and make appropriate business strategies. The file server receives data files periodically after every 3 hours.

Problem Statement:

- Determine top 10 station_id(s) where maximum number of songs were played, which were liked by unique users. It is not only the data which is important, rather it is the insight it can be used to generate important. Once we have made the data ready for analysis, we have to perform below analysis on a daily basis.
- Determine total duration of songs played by each type of user, where type of user can be 'subscribed' or 'unsubscribed'. An unsubscribed user is the one whose record is either not present in Subscribed_users lookup table or has *subscription_end_date* earlier than the *timestamp* of the song played by him.
- Determine top 10 connected artists. Connected artists are those whose songs are most listened by the unique users who follow them.
- Determine top 10 songs who have generated the maximum revenue. Royalty applies to a song only if it was *liked* or was *completed successfully* or both.
- Determine top 10 unsubscribed users who listened to the songs for the longest duration.

Solution:

There are five phases in this project:

- Data generation for web and mobile sources
- Data population in the lookup tables
- Data formatting
- Data enrichment
- Data analysis

In the course of converting an existing Hadoop project into a Spark project, I have written code for data formatting, data enrichment and data analysis phases.

Here are the rules to follow for data ingestion and formatting:

- Data coming from mobile applications reside in /data/mob and has csv format.
- Data coming from web applications reside in /data/web and has xml format.
- All the timestamp fields in data coming from web application is of the format YYYY-MM-DD HH:MM:SS.
- All the timestamp fields in data coming from mobile application is a long integer interpreted as UNIX timestamps.
- Finally, all timestamps must have the format of a long integer to be interpreted as UNIX timestamps.

Let's have a glance at the source code written in Scala language:

DataFormatting.scala

```
// import required Spark packages
import org.apache.spark.{SparkConf, SparkContext}
import org.apache.spark.sql

object DataFormatting {
  def main(args: Array[String]): Unit = {
    val conf = new SparkConf().setAppName("Data Formatting")
    val sc = new SparkContext(conf)           // create spark context using spark 'conf' object
    val sqlContext = new org.apache.spark.sql.hive.HiveContext(sc) //create a Hive context object
    val batchId = args(0)                     // get batch id from command line arguments
    // SQL query to create table 'formatted_input' with the given schema description
    val create_hive_table = s"""CREATE TABLE IF NOT EXISTS project.formatted_input
      (
        User_id STRING,
        Song_id STRING,
```

```

Artist_id STRING,
Timestamp STRING,
Start_ts STRING,
End_ts STRING,
Geo_cd STRING,
Station_id STRING,
Song_end_type INT,
Like INT,
Dislike INT
)
PARTITIONED BY
(batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
""

```

// Hive command to load **mobile data** from local file system to Hive table 'formatted_input'

```

val load_mob_data = s""""LOAD DATA LOCAL INPATH '/home/acadgild/project/data/mob/
file.txt' INTO TABLE project.formatted_input PARTITION (batchid='$batchId')""""

```

// Hive query to insert **web data** to Hive table 'formatted_input' converting time stamp format

```

val load_web_data = s""""INSERT INTO project.formatted_input
PARTITION(batchid='$batchId')
SELECT user_id,
song_id,
artist_id,
unix_timestamp(timestamp,'yyyy-MM-dd HH:mm:ss') AS timestamp,
unix_timestamp(start_ts,'yyyy-MM-dd HH:mm:ss') AS start_ts,
unix_timestamp(end_ts,'yyyy-MM-dd HH:mm:ss') AS end_ts,
geo_cd,
station_id,

```

```

        song_end_type,
        like,
        dislike
    FROM web_data
    """

    try {
// read web data stored in XML format from local file system
        val xmlData = sqlContext.read.format("com.databricks.spark.xml").option("rowTag",
"record").load("/home/acadgild/project/data/web/file.xml")
xmlData.createOrReplaceTempView("web_data")           //create a view from XML data loaded

        sqlContext.sql(create_hive_table)           // execute the HiveQL queries prepared above
        sqlContext.sql(load_mob_data)
        sqlContext.sql(load_web_data)
    }
    catch{
        case e: Exception=>e.printStackTrace()           // print stack trace in case of errors
    }
}

```

Rules for data enrichment:

- If any of like or dislike is NULL or absent, consider it as 0.
- If fields like Geo_cd and Artist_id are NULL or absent, consult the lookup tables for fields Station_id and Song_id respectively to get the values of Geo_cd and Artist_id.
- If corresponding lookup entry is not found, consider that record to be invalid.

DataEnrichment.scala

```

// import required Spark packages
import org.apache.spark.{SparkConf, SparkContext}
import org.apache.spark.sql

```

```

object DataEnrichment {
  def main(args: Array[String]): Unit = {
    val conf = new SparkConf().setAppName("Data Formatting")
    val sc = new SparkContext(conf)           //create spark context using spark 'conf' object
    val sqlContext = new org.apache.spark.sql.hive.HiveContext(sc) //create a Hive context object
    val batchId = args(0)                     // get batch id from command line arguments
    // Hive query to insert web data to Hive table 'enriched_data'
    val create_hive_table = s"""CREATE TABLE IF NOT EXISTS enriched_data
      (
        User_id STRING,
        Song_id STRING,
        Artist_id STRING,
        Timestamp STRING,
        Start_ts STRING,
        End_ts STRING,
        Geo_cd STRING,
        Station_id STRING,
        Song_end_type INT,
        Like INT,
        Dislike INT
      )
      PARTITIONED BY
      (batchid INT,
        status STRING)
      STORED AS ORC
    """

    // Hive query to insert data from 'formatted_output' into the table 'enriched_data'
    val load_data = s"""INSERT OVERWRITE TABLE enriched_data
      PARTITION (batchid, status)
      SELECT

```

```

i.user_id,
i.song_id,
sa.artist_id,
i.timestamp,
i.start_ts,
i.end_ts,
sg.geo_cd,
i.station_id,
IF (i.song_end_type IS NULL, 3, i.song_end_type) AS song_end_type,
IF (i.like IS NULL, 0, i.like) AS like,
IF (i.dislike IS NULL, 0, i.dislike) AS dislike,
i.batchid,
IF((i.like=1 AND i.dislike=1)
OR i.user_id IS NULL
OR i.song_id IS NULL
OR i.timestamp IS NULL
OR i.start_ts IS NULL
OR i.end_ts IS NULL
OR i.geo_cd IS NULL
OR i.user_id=""
OR i.song_id=""
OR i.timestamp=""
OR i.start_ts=""
OR i.end_ts=""
OR i.geo_cd=""
OR sg.geo_cd IS NULL
OR sg.geo_cd=""
OR sa.artist_id IS NULL
OR sa.artist_id=", 'fail', 'pass') AS status
FROM formatted_input i

```

```

LEFT OUTER JOIN station_geo_map sg ON i.station_id = sg.station_id
LEFT OUTER JOIN song_artist_map sa ON i.song_id = sa.song_id
WHERE i.batchid=$batchId
""""

try {
    sqlContext.sql("SET hive.auto.convert.join=false")           // configure Hive properties
    sqlContext.sql("SET hive.exec.dynamic.partition.mode=nonstrict")
    sqlContext.sql("USE project")                                // switch to an existing database 'project'

    sqlContext.sql(create_hive_table)                            // execute the HiveQL queries prepared above
    sqlContext.sql(load_data)
}
catch{
    case e: Exception=>e.printStackTrace()                       // print stack trace in case of errors
}
}
}

```

Data Analysis:

The following code contains HiveQL queries to obtain the data specified in the problem statement:

DataAnalysis.scala

```

// import required Spark packages
import org.apache.spark.{SparkConf, SparkContext}
import org.apache.spark.sql

object DataAnalysis {
    def main(args: Array[String]): Unit = {
        val conf = new SparkConf().setAppName("Data Analysis")
        val sc = new SparkContext(conf)                        //create spark context using spark 'conf' object
    }
}

```

```

val sqlContext = new org.apache.spark.sql.hive.HiveContext(sc) //create a Hive context object
val batchId = args(0) // get batch id from command line arguments

// Hive query to create a table to store top 10 station id's where maximum number of songs played
val create_top_10_stations = """CREATE TABLE IF NOT EXISTS top_10_stations
(
    station_id STRING,
    total_distinct_songs_played INT,
    distinct_user_count INT
)
PARTITIONED BY (batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE"""

// Hive query to determine top 10 station_id(s) and store results into 'top_10_stations' table
val load_top_10_stations = s"""INSERT OVERWRITE TABLE top_10_stations
PARTITION(batchid='$batchId')
SELECT station_id,
COUNT(DISTINCT song_id) AS total_distinct_songs_played,
COUNT(DISTINCT user_id) AS distinct_user_count
FROM enriched_data
WHERE status='pass'
AND batchid='$batchId'
AND like=1
GROUP BY station_id
ORDER BY total_distinct_songs_played DESC
LIMIT 10"""

```


// Hive query to create a table to store the total duration of songs played by each type of user

```
val create_users_behaviour = """CREATE TABLE IF NOT EXISTS users_behaviour
```

```
(
  user_type STRING,
  duration INT
)
PARTITIONED BY (batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE"""
```

// Hive query to determine total duration of songs played by each type of user and store results into

// 'users_behaviour' table

```
val load_users_behaviour = s"""INSERT OVERWRITE TABLE users_behaviour
```

```
  PARTITION(batchid='$batchId')
```

```
  SELECT
```

```
    CASE WHEN (su.user_id IS NULL OR CAST(ed.timestamp AS
    DECIMAL(20,0)) > CAST(su.subscn_end_dt AS
    DECIMAL(20,0))) THEN 'UNSUBSCRIBED'
```

```
    WHEN (su.user_id IS NOT NULL AND CAST(ed.timestamp AS
    DECIMAL(20,0)) <= CAST(su.subscn_end_dt AS
    DECIMAL(20,0))) THEN 'SUBSCRIBED'
```

```
  END AS user_type,
```

```
  SUM(ABS(CAST(ed.end_ts AS DECIMAL(20,0))-
  CAST(ed.start_ts AS DECIMAL(20,0)))) AS duration
```

```
  FROM enriched_data ed
```

```
  LEFT OUTER JOIN subscribed_users su
```

```
  ON ed.user_id=su.user_id
```

```
  WHERE ed.status='pass'
```

```
  AND ed.batchid='$batchId'
```

```
  GROUP BY CASE WHEN (su.user_id IS NULL OR
  CAST(ed.timestamp AS DECIMAL(20,0)) >
```

```
CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN  
'UNSUBSCRIBED'
```

```
WHEN (su.user_id IS NOT NULL AND CAST(ed.timestamp AS  
DECIMAL(20,0)) <= CAST(su.subscn_end_dt AS  
DECIMAL(20,0))) THEN 'SUBSCRIBED' END""
```

// Hive query to create a table to store data of top 10 connected artists.

```
val create_connected_artists = """"CREATE TABLE IF NOT EXISTS connected_artists
```

```
(  
  artist_id STRING,  
  user_count INT  
)  
PARTITIONED BY (batchid INT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE"""
```

// Hive query to determine top 10 connected artists and store results into 'connected_artists' table

```
val load_connected_artists = s""""INSERT OVERWRITE TABLE connected_artists
```

```
  PARTITION(batchid='$batchId')  
  SELECT ua.artist_id,  
  COUNT(DISTINCT ua.user_id) AS user_count  
  FROM  
  (  
    SELECT user_id, artist_id FROM users_artists  
    LATERAL VIEW explode(artists_array) artists AS artist_id  
  ) ua  
  INNER JOIN  
  (  
    SELECT artist_id, song_id, user_id  
    FROM enriched_data  
    WHERE status='pass'  
    AND batchid='$batchId'
```

```
) ed
ON ua.artist_id=ed.artist_id
AND ua.user_id=ed.user_id
GROUP BY ua.artist_id
ORDER BY user_count DESC
LIMIT 10"""
```

// Hive query to create a table to store data of top 10 royalty songs

```
val create_top_10_royalty_songs = """"CREATE TABLE IF NOT EXISTS top_10_royalty_songs
(
    song_id STRING,
    duration INT
)
PARTITIONED BY (batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE"""
```

// Hive query to determine top 10 royalty songs and store results into top_10_royalty_songs table

```
val load_top_10_royalty_songs = s""""INSERT OVERWRITE TABLE top_10_royalty_songs
PARTITION(batchid='$batchId')
SELECT song_id,
SUM(ABS(CAST(end_ts AS DECIMAL(20,0))-CAST(start_ts AS
DECIMAL(20,0)))) AS duration
FROM enriched_data
WHERE status='pass'
AND batchid='$batchId'
AND (like=1 OR song_end_type=0)
GROUP BY song_id
ORDER BY duration DESC
LIMIT 10"""
```

// Hive query to create a table to store data of top 10 unsubscribed users

```
val create_top_10_unsubscribed_users = """CREATE TABLE IF NOT EXISTS
top_10_unsubscribed_users

(
    user_id STRING,
    duration INT
)

PARTITIONED BY (batchid INT)

ROW FORMAT DELIMITED

FIELDS TERMINATED BY ','

STORED AS TEXTFILE"""
```

// Hive query to determine top 10 unsubscribed users who listened to songs for the longest duration

```
val load_top_10_unsubscribed_users = s"""INSERT OVERWRITE TABLE
top_10_unsubscribed_users

PARTITION(batchid='$batchId')

SELECT ed.user_id,

SUM(ABS(CAST(ed.end_ts AS DECIMAL(20,0))-
CAST(ed.start_ts AS DECIMAL(20,0)))) AS duration

FROM enriched_data ed

LEFT OUTER JOIN subscribed_users su

ON ed.user_id=su.user_id

WHERE ed.status='pass'

AND ed.batchid='$batchId'

AND (su.user_id IS NULL OR (CAST(ed.timestamp AS
DECIMAL(20,0)) > CAST(su.subscn_end_dt AS
DECIMAL(20,0))))

GROUP BY ed.user_id

ORDER BY duration DESC

LIMIT 10"""
```

```
try {
```

```
    sqlContext.sql("SET hive.auto.convert.join=false") // configure Hive properties
```

```

sqlContext.sql("USE project")           // switch to the working database 'project'
sqlContext.sql(create_top_10_stations)  // execute all the Hive queries created above
sqlContext.sql(load_top_10_stations)
sqlContext.sql(create_users_behaviour)
sqlContext.sql(load_users_behaviour)
sqlContext.sql(create_connected_artists)
sqlContext.sql(load_connected_artists)
sqlContext.sql(create_top_10_royalty_songs)
sqlContext.sql(load_top_10_royalty_songs)
sqlContext.sql(create_top_10_unsubscribed_users)
sqlContext.sql(load_top_10_unsubscribed_users)
}
catch{
  case e: Exception=>e.printStackTrace() // print stack trace in case of errors
}
}
}

```

Output:

Data Formatting – Screenshots of execution:

```
Applications Places System Mon Mar 19, 4:22 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
Table project.users_artists stats: [numFiles=1, numRows=0, totalSize=240, rawDataSize=0]
OK
Time taken: 2.868 seconds
[acadgild@localhost scripts]$ ls
create_hive_hbase_lookup.hql      data_enrichment.sh      generate_mob_data.py      populate-lookup.sh
create_hive_hbase_lookup.sh      data_export.sh          generate_mob_data.txt     start-daemons.sh
create_schema.sql                dataformatting.sh       generate_web_data.py      user-artist.hql
data_analysis.sh                 file_generation_in_python.txt hdfs:                    wrapper.sh
data_enrichment_filtering_schema.sh formatted_hive_load.hql  MusicDataAnalysis
[acadgild@localhost scripts]$ sh data_enrichment_filtering_schema.sh

Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-0.14.0.jar!/hive-log4j.properties
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/hive-jdbc-0.14.0-standalone.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
OK
Time taken: 1.374 seconds
OK
Time taken: 0.653 seconds
OK
Time taken: 0.066 seconds
OK
Time taken: 0.051 seconds
[acadgild@localhost scripts]$ sh data_formatting.sh
sh: data_formatting.sh: No such file or directory
[acadgild@localhost scripts]$ sh dataformatting.sh
Ivy Default Cache set to: /home/acadgild/.ivy2/cache
The jars for the packages stored in: /home/acadgild/.ivy2/jars
:: loading settings :: url = jar:file:/usr/local/spark/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml
com.databricks#spark-xml_2.10 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent;1.0
  confs: [default]
  found com.databricks#spark-xml_2.10;0.4.1 in central
downloading https://repo1.maven.org/maven2/com/databricks/spark-xml_2.10/0.4.1/spark-xml_2.10-0.4.1.jar ...
[acadgild@localhost:~/... acadgild@localhost:u...
```

```
Applications Places System Mon Mar 19, 4:44 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
.3 MB)
18/03/19 04:44:02 INFO SparkContext: Created broadcast 0 from newAPIHadoopFile at XmlFile.scala:46
18/03/19 04:44:05 INFO FileInputFormat: Total input paths to process : 1
18/03/19 04:44:06 INFO SparkContext: Starting job: treeAggregate at InferSchema.scala:109
18/03/19 04:44:06 INFO DAGScheduler: Got job 0 (treeAggregate at InferSchema.scala:109) with 1 output partitions
18/03/19 04:44:06 INFO DAGScheduler: Final stage: ResultStage 0 (treeAggregate at InferSchema.scala:109)
18/03/19 04:44:06 INFO DAGScheduler: Parents of final stage: List()
18/03/19 04:44:06 INFO DAGScheduler: Missing parents: List()
18/03/19 04:44:06 INFO DAGScheduler: Submitting ResultStage 0 (MapPartitionsRDD[3] at treeAggregate at InferSchema.scala:109)
, which has no missing parents
18/03/19 04:44:07 INFO MemoryStore: Block broadcast_1 stored as values in memory (estimated size 4.4 KB, free 366.1 MB)
18/03/19 04:44:07 INFO MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 2.6 KB, free 366.1 MB)
18/03/19 04:44:07 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on 192.168.43.118:38581 (size: 2.6 KB, free: 366.1 MB)
18/03/19 04:44:07 INFO SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1012
18/03/19 04:44:07 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 0 (MapPartitionsRDD[3] at treeAggregate at InferSchema.scala:109)
18/03/19 04:44:07 INFO TaskSchedulerImpl: Adding task set 0.0 with 1 tasks
18/03/19 04:44:08 INFO TaskSetManager: Starting task 0.0 in stage 0.0 (TID 0, localhost, partition 0, PROCESS_LOCAL, 5588 bytes)
18/03/19 04:44:09 INFO Executor: Running task 0.0 in stage 0.0 (TID 0)
18/03/19 04:44:09 INFO Executor: Fetching spark://192.168.43.118:50116/jars/com.databricks_spark-xml_2.10-0.4.1.jar with time stamp 1521414814465
18/03/19 04:44:10 INFO TransportClientFactory: Successfully created connection to /192.168.43.118:50116 after 779 ms (0 ms spent in bootstraps)
18/03/19 04:44:10 INFO Utils: Fetching spark://192.168.43.118:50116/jars/com.databricks_spark-xml_2.10-0.4.1.jar to /tmp/spark-92fa3996-1c5c-473c-b1e3-f552925e3bfb/userFiles-0d6b408c-467b-40c7-b5d0-ada8a49dc582/fetchFileTemp8207616526617628326.tmp
18/03/19 04:44:16 INFO Executor: Adding file:/tmp/spark-92fa3996-1c5c-473c-b1e3-f552925e3bfb/userFiles-0d6b408c-467b-40c7-b5d0-ada8a49dc582/com.databricks_spark-xml_2.10-0.4.1.jar to class loader
18/03/19 04:44:16 INFO Executor: Fetching spark://192.168.43.118:50116/jars/musicdataanalysis_2.11-1.0.jar with timestamp 1521414814470
18/03/19 04:44:16 INFO Utils: Fetching spark://192.168.43.118:50116/jars/musicdataanalysis_2.11-1.0.jar to /tmp/spark-92fa3996-1c5c-473c-b1e3-f552925e3bfb/userFiles-0d6b408c-467b-40c7-b5d0-ada8a49dc582/fetchFileTemp7398493232539183603.tmp
18/03/19 04:44:16 INFO Executor: Adding file:/tmp/spark-92fa3996-1c5c-473c-b1e3-f552925e3bfb/userFiles-0d6b408c-467b-40c7-b5d0-ada8a49dc582/musicdataanalysis_2.11-1.0.jar to class loader
18/03/19 04:44:18 INFO NewHadoopRDD: Input split: file:/home/acadgild/project/data/web/file.xml:0+6718
18/03/19 04:44:30 INFO Executor: Finished task 0.0 in stage 0.0 (TID 0). 2229 bytes result sent to driver
18/03/19 04:44:31 INFO TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 23001 ms on localhost (1/1)
```

acadgild@localhost:~/project/scripts

File Edit View Search Terminal Help

```
18/03/19 05:53:17 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create_database: Database(name=default, description:default database, locationUri:file:/home/acadgild/project/scripts/spark-warehouse, parameters:{})
18/03/19 05:53:18 INFO SparkSqlParser: Parsing command: web_data
```

```
18/03/19 05:53:19 INFO SparkSqlParser: Parsing command: CREATE TABLE IF NOT EXISTS project.formatted_input
```

```
(
  User_id STRING,
  Song_id STRING,
  Artist_id STRING,
  Timestamp STRING,
  Start_ts STRING,
  End_ts STRING,
  Geo_cd STRING,
  Station_id STRING,
  Song_end_type INT,
  Like INT,
  Dislike INT
)
PARTITIONED BY
(batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
```

```
18/03/19 05:53:19 INFO HiveMetaStore: 0: get_database: project
18/03/19 05:53:19 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 05:53:19 INFO HiveMetaStore: 0: get_database: project
18/03/19 05:53:19 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 05:53:20 INFO HiveMetaStore: 0: create_table: Table(tableName=formatted_input, dbName=project, owner=acadgild, createTime:1521418999, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name=user_id, type:string, comment:null), FieldSchema(name=song_id, type:string, comment:null), FieldSchema(name=artist_id, type:string, comment:null), FieldSchema(name=timestamp, type:string, comment:null), FieldSchema(name=start_ts, type:string, comment:null), FieldSchema(name=end_ts, type:string, comment:null), FieldSchema(name=geo_cd, type:string, comment:null), FieldSchema(name=station_id, type:string, comment:null), FieldSchema(name=song_end_type, type:int, comment:null), FieldSchema(name=like, type:int, comment:null), FieldSchema(name=dislike, type:int, comment:null)], location:null, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{field.delim=,, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocations:[])), partitionKeys:[FieldSchema(name=batchid, type:int, comment:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
```

acadgild@localhost:~/... acadgild@localhost:~/... Google Keep - Mozilla ...

acadgild@localhost:~/project/scripts

File Edit View Search Terminal Help

```
18/03/19 05:53:21 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=formatted_input
18/03/19 05:53:21 INFO HiveMetaStore: 0: get_table : db=project tbl=formatted_input
18/03/19 05:53:21 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=formatted_input
18/03/19 05:53:21 INFO HiveMetaStore: 0: get_partition_with_auth : db=project tbl=formatted_input[2]
18/03/19 05:53:21 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_partition_with_auth : db=project tbl=formatted_input[2]
```

```
18/03/19 05:53:23 INFO FileUtils: Creating directory if it doesn't exist: hdfs://localhost:9000/user/hive/warehouse/project.db/formatted_input/batchid=2
```

```
18/03/19 05:53:23 INFO SessionState: Could not get hdfsEncryptionShim, it is only applicable to hdfs filesystem.
```

```
18/03/19 05:53:25 INFO Hive: Renaming src: file:/home/acadgild/project/data/mob/file.txt, dest: hdfs://localhost:9000/user/hive/warehouse/project.db/formatted_input/batchid=2/file.txt, Status:true
```

```
18/03/19 05:53:26 INFO HiveMetaStore: 0: get_partition_with_auth : db=project tbl=formatted_input[2]
```

```
18/03/19 05:53:26 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_partition_with_auth : db=project tbl=formatted_input[2]
```

```
18/03/19 05:53:26 INFO HiveMetaStore: 0: append_partition : db=project tbl=formatted_input[2]
```

```
18/03/19 05:53:26 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=append_partition : db=project tbl=formatted_input[2]
```

```
18/03/19 05:53:26 WARN log: Updating partition stats fast for: formatted_input
```

```
18/03/19 05:53:26 WARN log: Updated size to 1238
```

```
18/03/19 05:53:26 INFO SparkSqlParser: Parsing command: INSERT INTO project.formatted_input
```

```
PARTITION(batchid='2')
SELECT user_id,
       song_id,
       artist_id,
       unix_timestamp(timestamp,'yyyy-MM-dd HH:mm:ss') AS timestamp,
       unix_timestamp(start_ts,'yyyy-MM-dd HH:mm:ss') AS start_ts,
       unix_timestamp(end_ts,'yyyy-MM-dd HH:mm:ss') AS end_ts,
       geo_cd,
       station_id,
       song_end_type,
       like,
       dislike
FROM web_data
```

```
18/03/19 05:53:26 INFO HiveMetaStore: 0: get_table : db=project tbl=formatted_input
```

```
18/03/19 05:53:26 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=formatted_input
```

```
18/03/19 05:53:27 INFO CatalystSqlParser: Parsing command: int
```

```
18/03/19 05:53:27 INFO CatalystSqlParser: Parsing command: string
```

```
18/03/19 05:53:27 INFO CatalystSqlParser: Parsing command: string
```

acadgild@localhost:~/... acadgild@localhost:~/... Google Keep - Mozilla ...

```
Applications Places System Mon Mar 19, 6:23 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
U117 S201 A300 1468094889 1494297562 1468094889 U ST401 1 0 1 2
Time taken: 1.309 seconds Fetched: 80 row(s)
hive> select * from formatted_input where batchid=2;
OK
U111 S205 A302 1495130523 1475130523 1485130523 AP ST413 1 0 1 2
U101 S205 A304 1475130523 1485130523 1485130523 A ST407 3 0 1 2
U116 S201 A305 1495130523 1465230523 1475130523 E ST410 1 0 0 2
U120 S207 A305 1465230523 1475130523 1465230523 U ST406 1 1 1 2
U114 S207 A303 1475130523 1485130523 1465230523 A ST401 2 1 1 2
S206 A304 1465230523 1485130523 1475130523 AP ST406 2 0 1 2
U117 S202 A302 1495130523 1475130523 1465130523 A ST406 2 0 1 2
U117 S209 A301 1465230523 1465130523 1465130523 AP ST406 1 1 1 2
U120 S210 A305 1465230523 1485130523 1465230523 ST413 3 0 0 2
U114 S208 1475130523 1465230523 1465230523 U ST413 0 1 0 2
U103 S200 A301 1465130523 1475130523 1465230523 U ST402 1 0 1 2
U104 S203 A305 1465230523 1465230523 1465230523 A ST401 2 0 0 2
U114 S210 A303 1465130523 1485130523 1475130523 U ST404 2 0 1 2
U116 S203 A302 1495130523 1465230523 1465230523 AP ST409 3 1 0 2
U100 S205 A301 1495130523 1465130523 1465230523 AP ST415 1 0 1 2
U105 S206 A300 1475130523 1475130523 1475130523 AP ST413 3 1 1 2
U111 S207 A304 1475130523 1485130523 1465130523 U ST405 2 1 1 2
U108 S203 A302 1475130523 1485130523 1465230523 A ST409 2 1 1 2
U107 S202 A305 1475130523 1475130523 1465130523 AU ST411 0 1 0 2
U117 S200 A303 1475130523 1465130523 1475130523 A ST400 3 0 1 2
U115 S209 A303 1494297562 1494297562 1462863262 A ST400 1 0 0 2
U109 S205 A302 1494297562 1462863262 1462863262 AP ST411 0 0 1 2
U109 S209 A301 1468094889 1465490556 1468094889 AP ST407 0 0 1 2
U108 S206 A303 1462863262 1468094889 1462863262 A ST407 2 1 0 2
U110 S206 A302 1468094889 1494297562 1494297562 E ST404 1 0 1 2
NULL S204 A301 1468094889 1465490556 1462863262 AP ST400 3 0 0 2
U112 S202 A303 1494297562 1468094889 1494297562 AU ST411 1 0 0 2
U112 S204 A305 1462863262 1494297562 1494297562 AP ST407 3 1 1 2
U115 S209 A300 1494297562 1468094889 1468094889 NULL ST410 0 1 1 2
U112 S208 NULL 1465490556 1462863262 1494297562 A ST408 2 1 0 2
U114 S208 A305 1465490556 1494297562 1468094889 A ST407 0 0 1 2
U112 S207 A301 1462863262 1462863262 1465490556 A ST400 2 0 0 2
U114 S207 A305 1494297562 1462863262 1465490556 A ST410 2 0 1 2
U114 S205 A304 1494297562 1462863262 1468094889 A ST414 3 0 1 2
U116 S204 A300 1494297562 1494297562 1462863262 U ST406 3 0 0 2
```

Data Enrichment – Screenshots of execution:

```
Applications Places System Mon Mar 19, 8:30 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
at org.apache.spark.deploy.SparkSubmit.main(SparkSubmit.scala)
data_enrichment.sh: line 11: --master: command not found
data_enrichment.sh: line 12: --jars: command not found
data_enrichment.sh: line 13: /home/acadgild/project/scripts/MusicDataAnalysis/target/scala-2.11/musicdataanalysis_2.11-1.0.jar: Permission denied
: command not found line 14:
data_enrichment.sh: line 33: syntax error: unexpected end of file
[acadgild@localhost scripts]$ vi data_enrichment.sh
[acadgild@localhost scripts]$ vi dataformatting.sh
[acadgild@localhost scripts]$ vi data_enrichment.sh
[acadgild@localhost scripts]$ sh data_enrichment.sh
data_enrichment.sh: line 3: unexpected EOF while looking for matching ``'
data_enrichment.sh: line 33: syntax error: unexpected end of file
[acadgild@localhost scripts]$ vi data_enrichment.sh
[acadgild@localhost scripts]$ sh data_enrichment.sh
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
18/03/19 08:30:10 INFO SparkContext: Running Spark version 2.0.0
18/03/19 08:30:14 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/03/19 08:30:15 WARN SparkConf:
SPARK_WORKER_INSTANCES was detected (set to '2').
This is deprecated in Spark 1.0+.

Please instead use:
- ./spark-submit with --num-executors to specify the number of executors
- Or set SPARK_EXECUTOR_INSTANCES
- spark.executor.instances to configure the number of instances in the spark config.

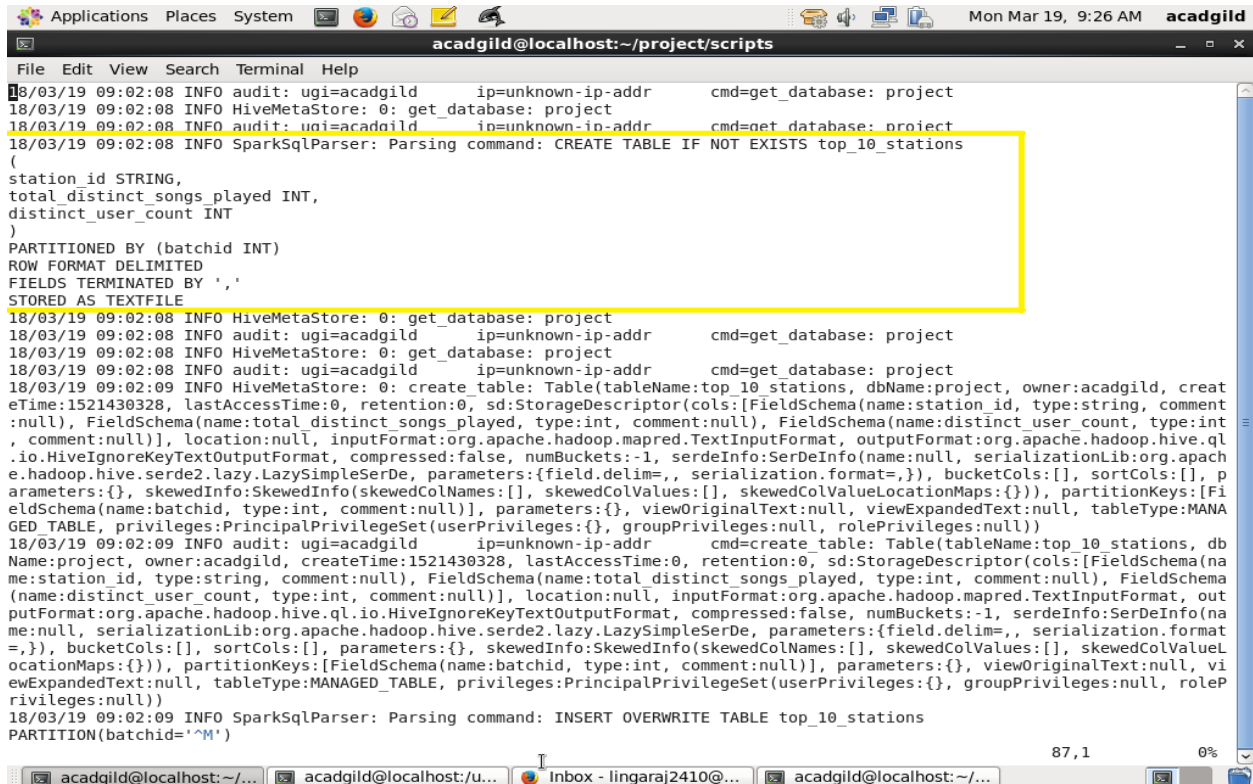
18/03/19 08:30:15 WARN Utils: Your hostname, localhost.localdomain resolves to a loopback address: 127.0.0.1; using 192.168.4.3.118 instead (on interface eth6)
18/03/19 08:30:15 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
18/03/19 08:30:16 INFO SecurityManager: Changing view acls to: acadgild
18/03/19 08:30:16 INFO SecurityManager: Changing modify acls to: acadgild
18/03/19 08:30:16 INFO SecurityManager: Changing view acls groups to:
18/03/19 08:30:16 INFO SecurityManager: Changing modify acls groups to:
18/03/19 08:30:16 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(acadgild); groups with view permissions: Set(); users with modify permissions: Set(acadgild); groups with modify permissions: Set()
```



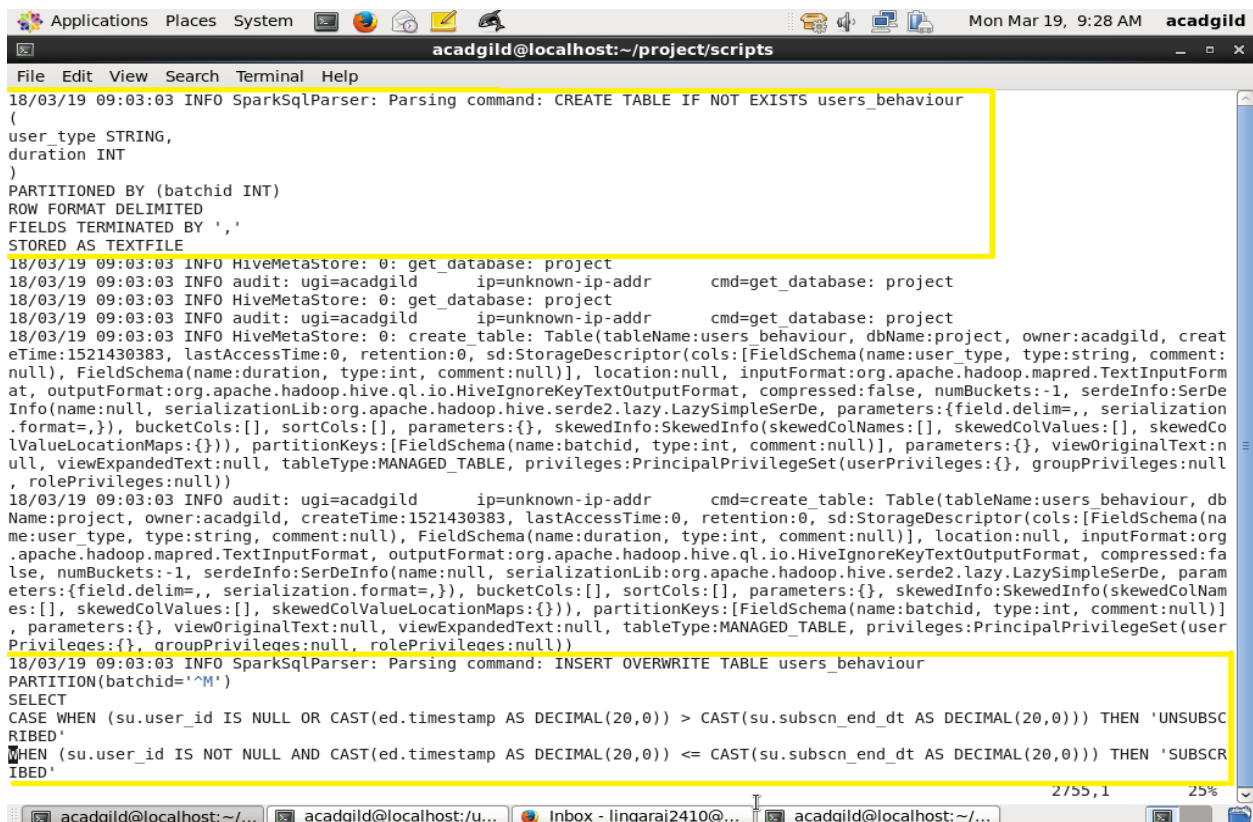
```
Applications Places System acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
6.2 MB)
18/03/19 08:33:19 INFO FileInputFormat: Total input paths to process : 2
18/03/19 08:33:19 INFO SparkContext: Starting job: sql at DataEnrichment.scala:73
18/03/19 08:33:19 INFO DAGScheduler: Got job 2 (sql at DataEnrichment.scala:73) with 2 output partitions
18/03/19 08:33:19 INFO DAGScheduler: Final stage: ResultStage 2 (sql at DataEnrichment.scala:73)
18/03/19 08:33:19 INFO DAGScheduler: Parents of final stage: List()
18/03/19 08:33:19 INFO DAGScheduler: Missing parents: List()
18/03/19 08:33:19 INFO DAGScheduler: Submitting ResultStage 2 (MapPartitionsRDD[19] at sql at DataEnrichment.scala:73), which
has no missing parents
18/03/19 08:33:19 INFO MemoryStore: Block broadcast_7 stored as values in memory (estimated size 96.9 KB, free 349.5 MB)
18/03/19 08:33:19 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 35.0 KB, free 349.5 MB)
18/03/19 08:33:19 INFO BlockManagerInfo: Added broadcast_7_piece0 in memory on 192.168.43.118:34164 (size: 35.0 KB, free: 366
.2 MB)
18/03/19 08:33:19 INFO SparkContext: Created broadcast 7 from broadcast at DAGScheduler.scala:1012
18/03/19 08:33:19 INFO DAGScheduler: Submitting 2 missing tasks from ResultStage 2 (MapPartitionsRDD[19] at sql at DataEnrichment.scala:73)
18/03/19 08:33:19 INFO TaskSchedulerImpl: Adding task set 2.0 with 2 tasks
18/03/19 08:33:19 INFO TaskSetManager: Starting task 0.0 in stage 2.0 (TID 2, localhost, partition 0, PROCESS_LOCAL, 6269 bytes)
18/03/19 08:33:19 INFO TaskSetManager: Starting task 1.0 in stage 2.0 (TID 3, localhost, partition 1, PROCESS_LOCAL, 6271 bytes)
18/03/19 08:33:19 INFO Executor: Running task 0.0 in stage 2.0 (TID 2)
18/03/19 08:33:19 INFO Executor: Running task 1.0 in stage 2.0 (TID 3)
18/03/19 08:33:19 INFO HadoopRDD: Input split: hdfs://localhost:9000/user/hive/warehouse/project.db/formatted_input/batchid=2/part-000000:0+1244
18/03/19 08:33:19 INFO HadoopRDD: Input split: hdfs://localhost:9000/user/hive/warehouse/project.db/formatted_input/batchid=2/file.txt:0+1238
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 105.48994 ms
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 14.468203 ms
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 35.611475 ms
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 76.418425 ms
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 42.880633 ms
18/03/19 08:33:20 INFO CodeGenerator: Code generated in 90.030894 ms
18/03/19 08:33:20 INFO LazyStruct: Missing fields! Expected 12 fields but only got 11! Ignoring similar problems.
18/03/19 08:33:20 INFO LazyStruct: Missing fields! Expected 12 fields but only got 11! Ignoring similar problems.
18/03/19 08:33:20 INFO SparkHiveDynamicPartitionWriterContainer: Sorting complete. Writing out partition files one at a time.
18/03/19 08:33:20 INFO SparkHiveDynamicPartitionWriterContainer: Sorting complete. Writing out partition files one at a time.
```

```
Applications Places System acadgild@localhost:usr/local/hive/lib
Time taken: 3.529 seconds, Fetched: 80 row(s)
hive> select * from enriched_data where batchid=2;
OK
U120 S207 A303 1465230523 1475130523 1465230523 AU ST406 1 1 1 2 fail
U114 S207 A303 1475130523 1485130523 1465230523 AU ST401 2 1 1 2 fail
S206 A302 1465230523 1485130523 1475130523 AU ST406 2 0 1 2 fail
U117 S209 A305 1465230523 1465130523 1465130523 AU ST406 1 1 1 2 fail
U120 S210 NULL 1465230523 1485130523 1465230523 J ST413 3 0 0 2 fail
U114 S210 NULL 1465130523 1485130523 1475130523 E ST404 2 0 1 2 fail
U100 S205 A301 1495130523 1465130523 1465230523 NULL ST415 1 0 1 2 fail
U105 S206 A302 1475130523 1475130523 1475130523 J ST413 3 1 1 2 fail
U111 S207 A303 1475130523 1485130523 1465130523 A ST405 2 1 1 2 fail
U108 S203 A303 1475130523 1485130523 1465230523 E ST409 2 1 1 2 fail
NULL S204 A304 1468094889 1465490556 1462863262 A ST400 3 0 0 2 fail
U112 S204 A304 1462863262 1494297562 1494297562 AP ST407 3 1 1 2 fail
U115 S209 A305 1494297562 1468094889 1468094889 A ST410 0 1 1 2 fail
U111 S205 A301 1495130523 1475130523 1485130523 J ST413 1 0 1 2 pass
U101 S205 A301 1475130523 1485130523 1485130523 AP ST407 3 0 1 2 pass
U116 S201 A301 1495130523 1465230523 1475130523 A ST410 1 0 0 2 pass
U117 S202 A302 1495130523 1475130523 1465130523 AU ST406 2 0 1 2 pass
U114 S208 A304 1475130523 1465230523 1465230523 J ST413 0 1 0 2 pass
U103 S200 A300 1465130523 1475130523 1465230523 AP ST402 1 0 1 2 pass
U104 S203 A303 1465230523 1465230523 1465230523 AU ST401 2 0 0 2 pass
U116 S203 A303 1495130523 1465230523 1465230523 E ST409 3 1 0 2 pass
U107 S202 A302 1475130523 1475130523 1465130523 A ST411 0 1 0 2 pass
U117 S200 A300 1475130523 1465130523 1475130523 A ST400 3 0 1 2 pass
U115 S209 A305 1494297562 1494297562 1462863262 A ST400 1 0 0 2 pass
U109 S205 A301 1494297562 1462863262 1462863262 A ST411 0 0 1 2 pass
U109 S209 A305 1468094889 1465490556 1468094889 AP ST407 0 0 1 2 pass
U108 S206 A302 1462863262 1468094889 1462863262 AP ST407 2 1 0 2 pass
U110 S206 A302 1468094889 1494297562 1494297562 E ST404 1 0 1 2 pass
U112 S202 A302 1494297562 1468094889 1494297562 A ST411 1 0 0 2 pass
U112 S208 A304 1465490556 1462863262 1494297562 E ST408 2 1 0 2 pass
U114 S208 A304 1465490556 1494297562 1468094889 AP ST407 0 0 1 2 pass
U112 S207 A303 1462863262 1462863262 1465490556 A ST400 2 0 0 2 pass
U114 S207 A303 1494297562 1462863262 1465490556 A ST410 2 0 1 2 pass
U114 S205 A301 1494297562 1462863262 1468094889 E ST414 3 0 1 2 pass
U116 S204 A304 1494297562 1494297562 1462863262 AU ST406 3 0 0 2 pass
U115 S206 A302 1465490556 1462863262 1462863262 AU ST401 2 1 0 2 pass
U110 S207 A303 1494297562 1468094889 1468094889 A ST411 3 1 0 2 pass
```

Data analysis – screenshots of execution:



```
Applications Places System Mon Mar 19, 9:26 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
18/03/19 09:02:08 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:02:08 INFO HiveMetaStore: 0: get_database: project
18/03/19 09:02:08 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:02:08 INFO SparkSqlParser: Parsing command: CREATE TABLE IF NOT EXISTS top_10_stations
(
  station_id STRING,
  total_distinct_songs_played INT,
  distinct_user_count INT
)
PARTITIONED BY (batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE
18/03/19 09:02:08 INFO HiveMetaStore: 0: get_database: project
18/03/19 09:02:08 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:02:08 INFO HiveMetaStore: 0: get_database: project
18/03/19 09:02:08 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:02:09 INFO HiveMetaStore: 0: create table: Table(tableName=top_10_stations, dbName=project, owner=acadgild, createTime:1521430328, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:station_id, type:string, comment:null), FieldSchema(name:total_distinct_songs_played, type:int, comment:null), FieldSchema(name:distinct_user_count, type:int, comment:null)], location:null, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{field.delim=, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocationMaps:{}), partitionKeys:[FieldSchema(name=batchid, type:int, comment:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:02:09 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create table: Table(tableName=top_10_stations, dbName=project, owner=acadgild, createTime:1521430328, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:station_id, type:string, comment:null), FieldSchema(name:total_distinct_songs_played, type:int, comment:null), FieldSchema(name:distinct_user_count, type:int, comment:null)], location:null, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{field.delim=, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocationMaps:{}), partitionKeys:[FieldSchema(name=batchid, type:int, comment:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:02:09 INFO SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE top_10_stations
PARTITION(batchid='^M')
87,1 0%
```



```
Applications Places System Mon Mar 19, 9:28 AM acadgild
acadgild@localhost:~/project/scripts
File Edit View Search Terminal Help
18/03/19 09:03:03 INFO SparkSqlParser: Parsing command: CREATE TABLE IF NOT EXISTS users_behaviour
(
  user_type STRING,
  duration INT
)
PARTITIONED BY (batchid INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE
18/03/19 09:03:03 INFO HiveMetaStore: 0: get_database: project
18/03/19 09:03:03 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:03:03 INFO HiveMetaStore: 0: get_database: project
18/03/19 09:03:03 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_database: project
18/03/19 09:03:03 INFO HiveMetaStore: 0: create table: Table(tableName=users behaviour, dbName=project, owner=acadgild, createTime:1521430383, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:user_type, type:string, comment:null), FieldSchema(name:duration, type:int, comment:null)], location:null, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{field.delim=, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocationMaps:{}), partitionKeys:[FieldSchema(name=batchid, type:int, comment:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:03:03 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create table: Table(tableName=users behaviour, dbName=project, owner=acadgild, createTime:1521430383, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:user_type, type:string, comment:null), FieldSchema(name:duration, type:int, comment:null)], location:null, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{field.delim=, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocationMaps:{}), partitionKeys:[FieldSchema(name=batchid, type:int, comment:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:03:03 INFO SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE users_behaviour
PARTITION(batchid='^M')
SELECT
CASE WHEN (su.user_id IS NULL OR CAST(ed.timestamp AS DECIMAL(20,0)) > CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'UNSUBSCRIBED'
WHEN (su.user_id IS NOT NULL AND CAST(ed.timestamp AS DECIMAL(20,0)) <= CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'SUBSCRIBED'
2/55,1 25%
```

Applications Places System Mon Mar 19, 9:30 AM acadgild

acadgild@localhost:~/project/scripts

```
File Edit View Search Terminal Help
18/03/19 09:05:47 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create table: Table(tableName:connected_artists,
dbName:project, owner:acadgild, createTime:1521430547, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(
name:artist_id, type:string, comment:null), FieldSchema(name:user_count, type:int, comment:null)], location:null, inputFormat
:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat, compresse
d:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, p
arameters:{field.delim=,, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedCo
lNames:[], skewedColValues:[], skewedColValueLocationMaps:{))), partitionKeys:[FieldSchema(name=batchid, type:int, comment:nu
ll)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(
userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:05:47 INFO SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE connected_artists
PARTITION(batchid='M')
SELECT
ua.artist_id,
COUNT(DISTINCT ua.user_id) AS user_count
FROM
(
SELECT user_id, artist_id FROM users_artists
LATERAL VIEW explode(artists_array) artists AS artist_id
) ua
INNER JOIN
(
SELECT artist_id, song_id, user_id
FROM enriched_data
WHERE status='pass'
AND batchid='M'
) ed
ON ua.artist_id=ed.artist_id
AND ua.user_id=ed.user_id
GROUP BY ua.artist_id
ORDER BY user_count DESC
LIMIT 10
18/03/19 09:05:48 INFO HiveMetaStore: 0: get_table : db=project tbl=users_artists
18/03/19 09:05:48 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=users_artists
18/03/19 09:05:48 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:05:48 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
@
4975,1 47%
```

acadgild@localhost:~/... acadgild@localhost:/u... Inbox - lingaraj2410@... acadgild@localhost:~/...

Applications Places System Mon Mar 19, 9:30 AM acadgild

acadgild@localhost:~/project/scripts

```
File Edit View Search Terminal Help
18/03/19 09:06:23 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create table: Table(tableName:top_10_royalty_son
s, dbName:project, owner:acadgild, createTime:1521430583, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSche
ma(name:song_id, type:string, comment:null), FieldSchema(name:duration, type:int, comment:null)], location:null, inputFormat
:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat, compressed
:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, p
arameters:{field.delim=,, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedCol
Names:[], skewedColValues:[], skewedColValueLocationMaps:{))), partitionKeys:[FieldSchema(name=batchid, type:int, comment:nu
ll)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilegeSet(u
serPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:06:23 INFO SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE top_10_royalty_songs
PARTITION(batchid='M')
SELECT song_id,
SUM(ABS(CAST(end_ts AS DECIMAL(20,0))-CAST(start_ts AS DECIMAL(20,0)))) AS duration
FROM enriched_data
WHERE status='pass'
AND batchid='M'
AND (like=1 OR song_end_type=0)
GROUP BY song_id
ORDER BY duration DESC
LIMIT 10
18/03/19 09:06:23 INFO HiveMetaStore: 0: get_table : db=project tbl=enriched_data
18/03/19 09:06:23 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=enriched_data
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:23 INFO CatalystSqlParser: Parsing command: int
18/03/19 09:06:23 INFO CatalystSqlParser: Parsing command: string
18/03/19 09:06:23 INFO CatalystSqlParser: Parsing command: string
18/03/19 09:06:23 INFO CatalystSqlParser: Parsing command: string
@
7660,1 72%
```

acadgild@localhost:~/... acadgild@localhost:/u... Inbox - lingaraj2410@... acadgild@localhost:~/...

```

18/03/19 09:06:29 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=create table: Table(tableName=top_10_unsubscribed
_users, dbName=project, owner=acadgild, createTime:1521430589, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[Fiel
dSchema(name=user_id, type:string, comment:null), FieldSchema(name=duration, type:int, comment:null)), location:null, inputFo
rmat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive ql.io.HiveIgnoreKeyTextOutputFormat, compr
essed:false, numBuckets:-1, serdeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerD
e, parameters:{field.delim=, serialization.format=,}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skew
edColNames:[], skewedColValues:[], skewedColValueLocationMaps:{})), partitionKeys:[FieldSchema(name=batchid, type:int, commen
t:null)], parameters:{}, viewOriginalText:null, viewExpandedText:null, tableType:MANAGED_TABLE, privileges:PrincipalPrivilege
Set(userPrivileges:{}, groupPrivileges:null, rolePrivileges:null))
18/03/19 09:06:29 INFO SparkSqlParser: Parsing command: INSERT OVERWRITE TABLE top_10_unsubscribed_users
PARTITION(batchid='M')
SELECT
ed.user_id,
SUM(ABS(CAST(ed.end_ts AS DECIMAL(20,0))-CAST(ed.start_ts AS DECIMAL(20,0)))) AS duration
FROM enriched_data ed
LEFT OUTER JOIN subscribed_users su
ON ed.user_id=su.user_id
WHERE ed.status='pass'
AND ed.batchid='M'
AND (su.user_id IS NULL OR (CAST(ed.timestamp AS DECIMAL(20,0)) > CAST(su.subscn_end_dt AS DECIMAL(20,0))))
GROUP BY ed.user_id
ORDER BY duration DESC
LIMIT 10
18/03/19 09:06:29 INFO HiveMetaStore: 0: get_table : db=project tbl=enriched_data
18/03/19 09:06:29 INFO audit: ugi=acadgild ip=unknown-ip-addr cmd=get_table : db=project tbl=enriched_data
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO Persistence: Request to load fields "comment,name,type" of class org.apache.hadoop.hive.metastore.mode
l.MFieldSchema but object is embedded, so ignored
18/03/19 09:06:29 INFO CatalystSqlParser: Parsing command: int

```

```

hive> select * from top_10_stations where batchid=2;

```

```

OK
ST411 2 2 2
ST409 1 1 2
ST401 1 1 2
ST408 1 1 2
ST407 1 1 2
ST413 1 1 2
Time taken: 0.655 seconds, Fetched: 6 row(s)

```

```

hive>
hive> select * from users_behaviour where batchid=2;

```

```

OK
UNSUBSCRIBED 173032867 2
SUBSCRIBED 106807233 2
Time taken: 0.234 seconds, Fetched: 2 row(s)
hive>

```

```

hive> select * from connected_artists where batchid=2;

```

```

OK
A301 4 2
A302 3 2
A304 1 2
Time taken: 0.116 seconds, Fetched: 3 row(s)
hive>

```

```

hive> select * from top_10_royalty_songs where batchid=2;

```

```

OK
S208 57636973 2
S202 10000000 2
S206 5231627 2
S209 2604333 2
S207 0 2
S205 0 2
S203 0 2
Time taken: 0.143 seconds, Fetched: 7 row(s)

```

```

hive>
hive> select * from top_10_unsubscribed_users where batchid=2;

```

```

OK
U117 46202673 2
U116 41334300 2
U115 31434300 2
U112 26202673 2
U107 10000000 2
U111 10000000 2
U114 7858921 2
U109 0 2
U110 0 2
Time taken: 0.155 seconds, Fetched: 9 row(s)
hive>

```